
CSE517A – HOMEWORK 3

M. Neumann

Mar 27 2018

- Please keep your written answers brief and to the point. Incorrect or rambling statements can hurt your score on a question.
- If your hand writing is not readable, we **cannot give you credit**. We recommend you type your solutions in \LaTeX and compile a .pdf for each answer. **Start every problem on a new page!**
- This will be due THU **April 12 2018** at **10am** with an automatic 3-day extension.
- You may work in groups of at most 2 students.
- Submission instructions:
 - Start every problem on a **new page**.
 - Submissions will be exclusively accepted via **Gradescope**. Find instructions on how to get your Gradescope account and submit your work on the course webpage.

Problem 1 (30 points) Parameter Learning for Gaussian Processes (GPs)

For simplicity you may assume zero-mean observations for the entire problem.

- (a) (9 pts) Warm-up: assuming noise-free training data $D = \{(\mathbf{x}_i, f_i)\}_{i=1,\dots,n}$ with $f_i = f(\mathbf{x}_i)$, show that the variance cov_{f_i} for the GP prediction for a training point \mathbf{x}_i is 0.
- (b) (9 pts) Despite being a non-parametric model, we still have to learn the kernel parameters θ for a Gaussian process. Those so called *hyperparameters* can be learned by maximizing the probability of observing the training data given the GP prior. This can be formally expressed by the *marginal likelihood* $p(\mathbf{y} \mid X, \theta)$. Luckily for standard GPR, this marginal likelihood can be computed in closed form. Derive the **analytic log marginal likelihood expression** (assuming a GP prior, noisy observations with Gaussian i.i.d. noise $\epsilon \sim \mathcal{N}(0, \sigma_n)$, and a parameterized covariance/kernel function K_θ). By using K_θ we indicate that the kernel matrix K depends on the kernel parameters θ .

HINT: use the fact that $\mathbf{y} \sim \mathcal{N}(0, K_y)$, where $K_y = K_\theta + \sigma^2 I$.

- (c) (6 pts) When implementing GPs (prediction and learning method), we aim to compute the required inverse matrix K_y^{-1} as efficient as possible. To do so, we leverage the Choleskey decomposition of $K_y = LL^T$, where L is a lower triangular matrix. State the log marginal likelihood in terms of α and L , where $\alpha = L^T \backslash (L \backslash \mathbf{y})$ and \backslash denotes left-division indicating that we solve a system of linear equations $L\mathbf{b} = \mathbf{y}$ for \mathbf{b} . (i.e., $L \backslash \mathbf{y} = L^{-1}\mathbf{y} = \mathbf{b} \Leftrightarrow L\mathbf{b} = \mathbf{y}$).
- (d) (6 pts) For learning our goal is to pick the hyper-parameters θ that maximize the log marginal likelihood (or minimize the negative log marginal likelihood). Derive the **derivative of the log marginal likelihood**. Again, state this equation in terms of α and L . This expression is used in a gradient descent procedure in a practical implementation.

Problem 2 (25 points) k -means Clustering

- (a) (10 pts) Consider the following two possible termination conditions for k -means:
- (i) STOP if assignments do not change
 - (ii) STOP if cluster centers do not change
- Are these two conditions equivalent to each other? I.e. (i) \Leftrightarrow (ii)? Prove your answer.
- (b) (5 pts) Does the k -means algorithm always converge? Argue why or why not.
- (c) (5 pts) Is it possible that the k -means algorithm generates empty clusters? Argue why or why not.
- (d) (5 pts) Is it possible to find non-convex clusters by the k -means algorithm? Argue why or why not.

Problem 3 (35 points) Expectation-Maximization (EM) for Mixture Model Clustering

We are given a set of data points $X = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ with $\mathbf{x}_i \in \mathbb{R}^d$, which are assumed to be drawn from a probabilistic model consisting of k probability distributions:

$$g(\mathbf{x} \mid \Theta) = \sum_{j=1}^k \pi_j p(\mathbf{x} \mid \boldsymbol{\theta}_j)$$

where π_j is the probability of drawing from the j -th distribution (i.e. $\sum \pi_j = 1$), and Θ is our collection of parameters (i.e. $\Theta = \{(\pi_j, \boldsymbol{\theta}_j)\}_{j=1}^k$).

For a Gaussian mixture component:

$$p(\mathbf{x} \mid \boldsymbol{\theta}_j) \sim \mathcal{N}_d(\boldsymbol{\mu}_j, \Sigma_j)$$

we have $\boldsymbol{\theta}_j = \{\boldsymbol{\mu}_j, \Sigma_j\}$, where $\boldsymbol{\mu}_j \in \mathbb{R}^d$ and $\Sigma_j \in \mathbb{R}^{d \times d}$ are the mean and covariance of the multivariate normal distribution.

- (a) (5 pts) Warm-up: what is the probability that some point \mathbf{x} belongs to the j -th distribution for a *Gaussian mixture model*?
- (b) (5 pts) Our Expectation-Maximization (EM) clustering algorithm needs a criteria for convergence. One method would be to assess the *likelihood of the data* under our previously estimated parameters Θ and under our current parameters Θ' . With some chosen $\epsilon > 0$, if our updated parameters have a likelihood that is ϵ greater than the likelihood of the previous estimated parameters, we would replace the current estimates with our updated parameters and continue with our algorithm. Otherwise, we exit. State the likelihood $\mathcal{L}(X \mid \Theta)$, for the data X given the parameters Θ for a **general mixture model** g .
- (c) (10 pts) Write pseudocode for an algorithm to find satisfying parameter estimations for a **general mixture model** using $Z = [\mathbf{z}_1, \dots, \mathbf{z}_n]^\top$ to summarize the cluster membership probabilities for all data points \mathbf{x}_i and the termination criteria introduced in the previous part.
- (d) (15 pts) Despite it's intuitive interpretation, it turns out that expectation maximization performs MLE. More specifically, it maximizes the following lower bound of the likelihood:

$$\mathcal{B} = \sum_{i=1}^n \sum_{j=1}^k z_{ij} \log \left(\frac{\pi_j p(\mathbf{x}_i \mid \boldsymbol{\theta}_j)}{z_{ij}} \right),$$

where $\sum_{j=1}^k z_{ij} = 1$ and $z_{ij} = [\mathbf{z}_i]_j$.¹

Now, let's assume a mixture model for d -dimensional **binary input data**, where each mixture component is represented as a product of Bernoulli distributions:

$$p(\mathbf{x}_i \mid \boldsymbol{\theta}_j) = \prod_{m=1}^d (\theta_{jm})^{x_{im}} (1 - \theta_{jm})^{(1-x_{im})}$$

¹To be able to use this upper bound (also know as *Jensen's inequality*) we had to cast the likelihood as an expectation. That's why the z_{ij} 's appear in the equation.

with θ_j being the vector of dimension-specific probabilities for the j th mixture component. Start by writing down the log-likelihood for this specific distribution. Now, instead of maximizing this log-likelihood, we maximize the lower bound given as \mathcal{B} . Based on this maximization, derive the MLE estimate for the parameters θ_{jm} , which corresponds to the update rule used in the M-step in the EM algorithm.